



KubeCon



CloudNativeCon

North America 2019

Kubeflow: Multi-Tenant, Self-Serve, ML Platform

kunming@google.com
kam.d.kasravi@intel.com

Enabling the end user

- **Multi-tenant self-serve workspaces for developers and data scientists**
- **Do not saddle the end user with k8s details**
- **Deploy job to the right CPU/Accelerated hardware**
- **Kustomize overlays for different cpu/accelerated hardware combinations**



Kubeflow: A platform for building ML products

- **Leverage containers and Kubernetes to solve the challenges of building ML products**
 - Reduce the time and effort to get models launched
- **Why Kubernetes**
 - Kubernetes has won
 - Kubernetes runs everywhere
 - Enterprises can adopt shared infrastructure and patterns for ML and non ML services
 - Knowledge transfer across the organization
- **Kubeflow is open**
 - No lock in
 - 200 Members
 - 20+ [Organizations](#)
 - Stats available @ <http://devstats.kubeflow.org>



Kubeflow Cloud Providers

- Google Kubernetes Engine
- AWS
- Azure

Kubeflow Native K8

- Deployable to any k8 existing cluster



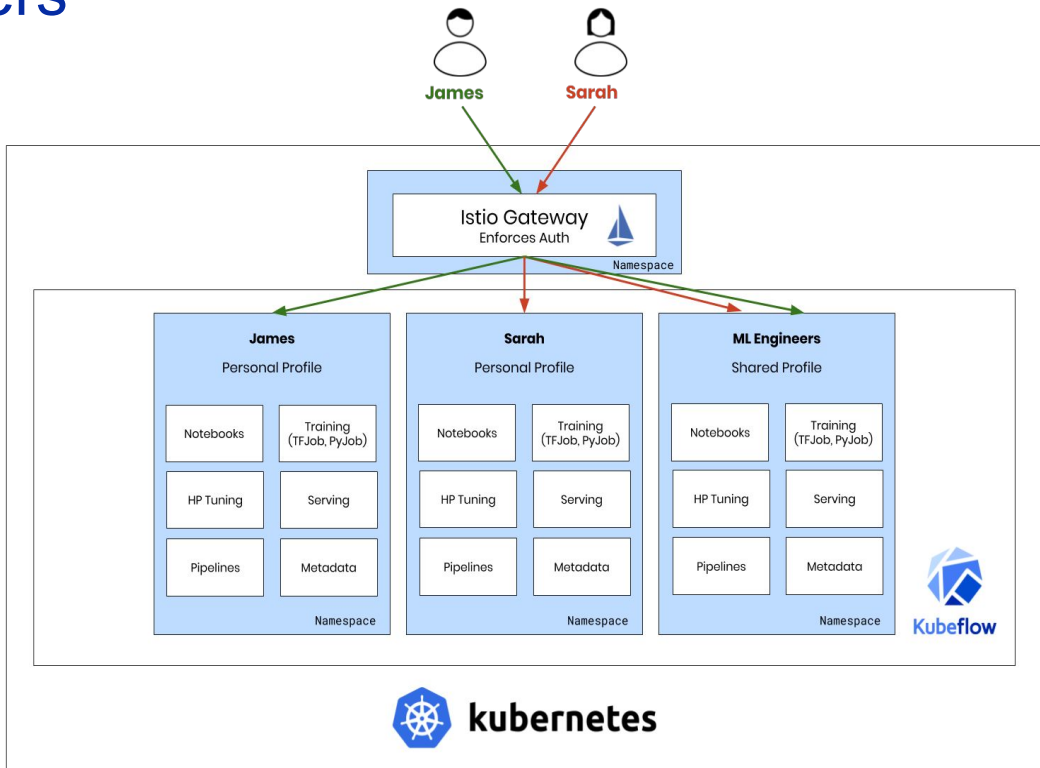
ML Applications

- Goal: applications for every stage of ML
- Examples:
 - Experimentation / Data Exploration
 - Jupyter/JupyterHub
 - Training
 - Tensorflow & Pytorch distributed training managed through K8s CRDs
 - Katib - HP Tuning
 - Workflows:
 - Pipelines
 - Metadata
 - Tracking and managing metadata of ML workflows
 - Feature Store
 - Feast (from GOJEK)



Multi Tenancy for End Users

Users will operate on same k8s cluster while each user has their own workspace hosting their services. Workspaces are logically isolated: each user can only access services to their own workspace.



K8s Multi-Tenancy Challenges

- Define clear user workspace boundary for access isolation
- K8s in-cluster network is transparent
 - Services are default visible from all pods
 - Need to establish network access control
- Access control around traffic through shared ingress
 - Users might access services in their own workspaces through same ingress.
 - Need to establish access control behind ingress: user can only access workspace after permission check
- Workspace access sharing & revoke
 - Each workspace owner should be able to share/revoke workspace access
 - Access sharing should not leak owner privilege while allow invited user operating on CRs
- All policies, roles and bindings involved should behave in consistency.



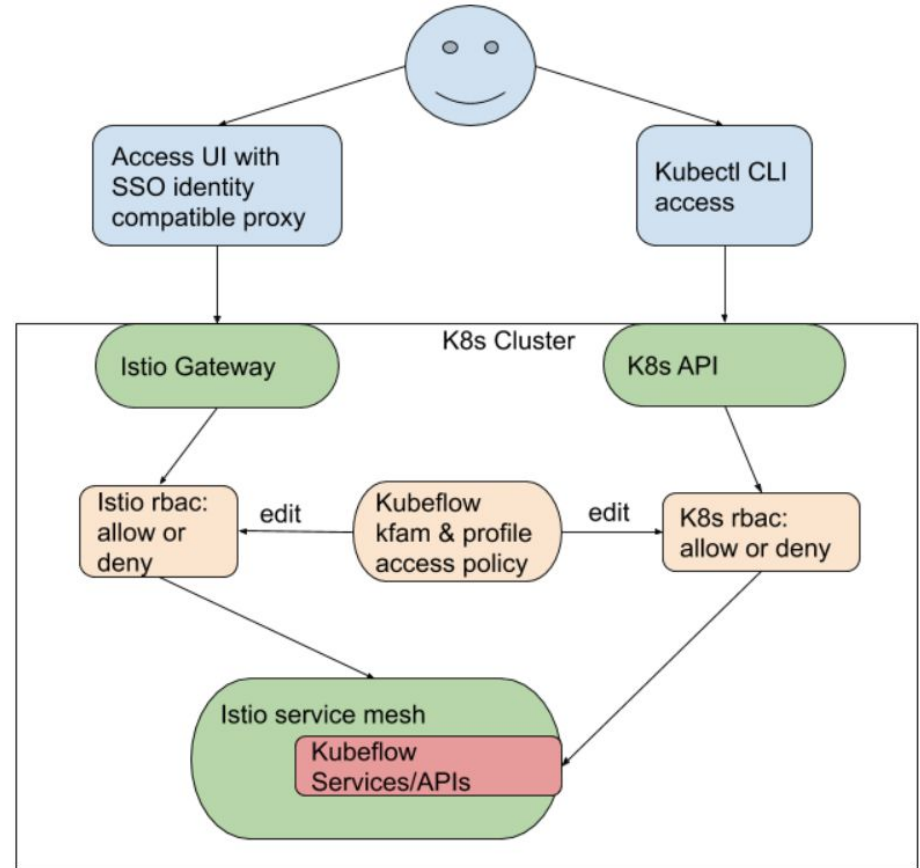
Kubeflow Multi-Tenancy (Profiles)

- Define user **workspace as namespace** and build access control around it
 - Manage user access to namespace through **k8s rbac policy**.
- Leverage Istio to **control in-cluster traffic**
 - By default requests to user workspaces are denied unless allowed by **Istio Rbac**
- Leverage **Identity-Aware Proxy** and Istio to control **traffic through ingress**
 - Identity user request through Identity-Aware Proxy.
 - Istio then do rbac check on request target workspace and identity
- Enable workspace access **sharing & revoke**
 - Workspace owners can share/revoke workspace access with other users **through kubeflow UI**
 - Invited users will have k8s edit permission plus permission to operate kubeflow CRs
- Self-serve
 - New user can self-register to create and own their workspace **through kubeflow UI**
- **Kubeflow Profile CR** to control all policies, roles and bindings involved and **guarantee consistency**.
 - Offer plugin interface to manage external resource/policy outside k8s, eg. access control of public cloud APIs



Kubeflow Access Control

- User access through kubectl: controlled by k8s rbac policy.
- User access through browser: controlled by istio rbac policy.
- Kubeflow multi-tenancy is implemented **k8s-native** way, new services can be integrated easily.



Kubeflow Profile

Created by the user via cli: *kubectl apply -f myprofile.yaml* or *kubeflow UI*

```
apiVersion: kubeflow.org/v1alpha1
kind: Profile
metadata:
  name: $(name)
spec:
  owner: $(owner)
```

Data scientists use Profiles to create various types of workspaces, where they can run training, inference, etc.



Create Kubeflow Profile

yaml

ui

```
apiVersion: kubeflow.org/v1beta1
kind: Profile
metadata:
  name: demo-namespace # profile name is also namespace name
spec:
  owner:
    kind: User
    name: user1@email.com # replace with the email of the user
  plugins:
    - kind: WorkloadIdentity
  spec:
    gcpServiceAccount: user1-gcp@project-id.iam.gserviceaccount.com
```



Namespace

A namespace is a collection of Kubeflow services. Resources created within a namespace are isolated to that namespace. By default, a namespace will be created for you.

Namespace Name

demo-namespace

Go back

Finish



Share Kubeflow Profile with other users

The screenshot shows the 'Manage Contributors' interface in Kubeflow. A red arrow labeled 'owner' points to the 'Account info (Cluster Admin)' section, which shows the user 'abhishek@google.com'. Another red arrow labeled 'contributor' points to the 'Contributors to your namespace - kubeflow-abhishek' section, which shows a table of contributors. The table has columns for 'Namespace', 'Owner', and 'Contributors'. The first row shows 'demo namespace' with 'kunming@google.com' as the owner and 'kunming@google.com' as the contributor. The second row shows 'kubeflow-abhishek' with 'abhishek@google.com' as the owner and 'kunming@google.com' as the contributor.

Manage Contributors

Account info (Cluster Admin)
abhishek@google.com

Namespace memberships

Namespaces	Your role
kubeflow-abhishek	Owner

Contributors to your namespace - kubeflow-abhishek

Email Addresses
kunming@google.com x Add by email address

Cluster namespaces

Namespace	Owner	Contributors
demo namespace	kunming@google.com	kunming@google.com
kubeflow-abhishek	abhishek@google.com	kunming@google.com



Kubeflow Device Overlays (accelerators, cpus)

- Device Overlays into Profiles, Pods
- Uses Profile extensions, Tekton Pipelines
- Can also be applied to Argo workflows and other pipeline engines



Kubeflow Profile adds cpu/accelerator quotas

Has a quotas section that is added to the namespace

```
apiVersion: kubeflow.org/v1alpha1
```

```
kind: Profile
```

```
metadata:
```

```
  name: $(name)
```

```
spec:
```

```
  owner: $(owner)
```

```
  quota:
```

```
    hard:
```

```
      requestsCpu: $(requestsCpu)
```

```
      requestsMemory: $(requestsMemory)
```

```
      requestsGpu: $(requestsGpu)
```

```
      limitsCpu: $(limitsCpu)
```

```
      limitsMemory: $(limitsMemory)
```

```
      <vendor/device>: <value>
```

This could be added by an admission-controller or gitops



Tekton Pipeline to run a model

kfctl can deploy manifest files from different repos or on disk

tk/tk-pipeline-run/overlays/cpu-node-selector

- kustomization.yaml
- params.env
- params.yaml
- **pipeline-run.yaml**



Tekton Pipeline to run a model

Adds a podTemplate.nodeSelector.cpu value

tk/tk-pipeline-run/overlays/cpu-node-selector/pipeline-run.yaml

```
apiVersion: tekton.dev/v1alpha1
kind: PipelineRun
metadata:
  name: $(generateName)
spec:
  podTemplate:
    nodeSelector:
      cpu: "$(cpuType)"
```



Tekton Pipeline to run a model

kfctl can deploy manifest files from different repos or on disk

tk/tk-pipeline/overlays/run-model-task

- config-map.yaml
- kustomization.yaml
- params.env
- params.yaml
- pipeline_patch.yaml
- **task.yaml**

tk/tk-pipeline/overlays/run-model-task-nvidia

- kustomization.yaml
- params.env
- params.yaml
- **task_patch.yaml**



Tekton Pipeline to run a model

run-model example

tk/tk-pipeline-run/overlays/run-model-task/task.yaml

```
apiVersion: tekton.dev/v1alpha1
kind: Task
metadata:
  name: run-model
spec:
  inputs:
    params:
      - name: imageName
        type: string
  steps:
    - name: run-model
      image: $(inputs.params.imageName)
      command: ["/bin/bash", "/run-model/run-model.sh"]
```



Tekton Pipeline to run a model

task is patched to add gpu info

tk/tk-pipeline-run/overlays/run-model-task-nvidia/task_patch.yaml

```
- op: add
  path: /spec/steps/0/resources
  value:
    limits:
      nvidia.com/gpu: $(accelerator_count)
```



Tekton Pipeline to run a model

config file to run a model on a node with cpu=skylake, nvidia.com/gpu

kfctl apply -f kfdef/run-model-gpu.yaml

```
- kustomizeConfig:
  overlays:
    - run-model-task
    - run-model-task-nvidia
  parameters:
    - name: accelerator_count
      value: 1
  repoRef:
    name: manifests
    path: tk/tk-pipeline
  name: tk-pipeline
```

```
- kustomizeConfig:
  overlays:
    - application
    - cpu-node-selector
  parameters:
    - name: cpuType
      value: skylake
  repoRef:
    name: manifests
    path: tk/tk-pipeline-run
  name: tk-pipeline-run
```



DEMO

- A kubeflow deployment that created a GKE cluster with 2 nodes

CPU Platform	Accelerator Type	Machine Type	Image Type
Intel Skylake	gpu nvidia-tesla-t4	n1-standard-8	cos
Intel Cascade Lake	-	c2-standard-8	ubuntu

- Run the same tensorflow model within a Profile but with different overlays

Pod limits selects the accelerator type (nvidia.com/gpu: '1')

Pod affinity selects the cpu platform (cpu: cascadelake)





Kubeflow

Thank You

- Kubeflow website - <https://www.kubeflow.org/>
- Code - <https://github.com/kubeflow/kubeflow>

